

UNIVERSITY OF WATERLOO
Faculty of Mathematics

OPTIMIZING PERFORMANCE ENGINEERING

Veeva Systems
Toronto, Ontario

Prepared by
Gary Jiayi Lin
2B Computer Science
ID 20557560
August 23rd, 2016

MEMORANDUM

To: Stefan Timofte

From: Gary Lin

Date: August 23rd, 2016

Re: Work Report: Optimizing Performance Engineering

As we agreed, I have prepared the enclosed report, “Optimizing Performance Engineering,” for my 2B work report and for the Veeva Toronto Research and Development Teams. This report, the second of four work reports that the Co-operative Education Program requires that I successfully complete as part of my BCS Co-op degree requirements, has not received academic credit.

The Vault Performance team that you lead provides systems performance monitoring, measurements and optimization for both internal software developers and external customers. My job as Intern Performance Engineer required that I develop and execute performance tests, develop data collection/organization scripts and provide overall support for your work. This report is an in-depth study of the optimal schedule for whom and when to develop/execute particular performance tests for a product such as Veeva Vault.

The Faculty of Mathematics requests that you evaluate this report for command of topic and technical content/analysis. Following your assessment, the report, together with your evaluation, will be submitted to the Math Undergrad Office for evaluation on campus by qualified work report markers. The combined marks determine whether the report will receive credit and whether it will be considered for an award.

Thank you for your assistance in preparing this report.

Gary Jiayi Lin

Table of Contents

List of Figures	ii
Executive Summary	iii
Introduction.....	1
Analysis.....	2
The Problem.....	2
Emphasis on Architectural Design.....	3
Emphasis on Agile Reinforcement	5
Performance Unit Testing	5
Performance Regression Testing	6
Disadvantages to Performance Investment	6
Conclusion	7
Recommendations	8
References	9

List of Figures

Figure 1: Diagram to show diminishing returns in performance as extra labor is employed	2
Figure 2: The Agile Development Methodology Lifecycle.....	4

Executive Summary

“Optimizing Performance Engineering” focuses on the key requirements in systems engineering for software performance. This report seeks to discover the optimal integrated schedule for activities and practices related to software performance in moderate to large scale software companies. This report aims to provide relief to the painful traditional practices of software performance engineering.

Specifically, this report focuses on early stage planning and system architectural design. It describes the inevitable causal effect a system’s architecture has on its performance in production, and how to ensure that it is carefully constructed.

Furthermore, “Optimizing Performance Engineering” critically analyzes the agile philosophy and the focus on adaptive development rather than sequential development, and produces improvements to the agile sprint loop. It is recommended that at the beginning of every sprint, sprint planning incorporates a minimum of 25% performance optimization work. The report also advises that developers write performance unit tests for their code, on top of the functional unit tests that they should be writing. Lastly, the report also details running performance regression tests at the end of every sprint to track the performance progress of the system.

This report also briefly depicts the issues that may arise from investing increased resources towards performance work. Mainly, there would be a functionality compromise, as less resources would be available for thorough functional testing.

In conclusion, “Optimizing Performance Engineering” localizes a complex study of software development and produces a set of measures that will ultimately lead to a streamlined software performance regulation routine, thus optimizing performance engineering.

Introduction

“Problems worthy of attack prove their worth by means of retaliation” stated Danish scientist, inventor and mathematician Piet Hein (Hein, 2). Amongst the thousands of problems software engineers face throughout their careers, software performance quality happens to be one of the most complex and difficult to solve. The struggle to meet software performance targets is comparable to one climbing a ski slope with their skis still on. Also, the target goal, a flag, is positioned halfway up the hill, not at the top. Getting to the flag is difficult enough, as the slipperiness ultimately leads to a long slide down to the bottom of the hill, along with endless frustration. Eventually when one finally makes it to flag, they would discover that they’ve already drained all their energy. Yet, due to the slipperiness and angle of the hill, one would have to use energy simply to maintain the altitude and avoid accidentally rolling to the bottom.

Considering the hill as a measure of software performance, the flag as a foreseeable performance target, the act of climbing the hill as an effort to reach the target, and the energy used as resources invested, the imagery sharpens. Software performance engineering is difficult. It is a problem undoubtedly worthy of attack, and it undoubtedly retaliates.

However, there exists ample tools, strategies, methodologies and practices to ensure that software performance targets are met on a regular basis. This report will evaluate the general problem of poor software performance, the effects this problem has on the software industry, various performance engineering practices and the optimal performance-focused development cycle. Perspectives in this report will alternate between reference to the general software industry and the specific work methodologies I grew accustomed to as a member of the Vault team at Veeva Systems.

Analysis

The Problem

When developing software, there really only are two questions of interest. Does this software do what it needs to do? If so, does it do it efficiently? Making software that does exactly what is needed is difficult enough, which is why the majority of efforts and resources are dedicated to ensuring that software simply functions as needed. Thus, often, little or no attention is paid to software performance. However, according to Dr. Williams and Dr. Smith from Performance Engineering Services, “poor performance costs the software industry millions of dollars annually in lost revenue, decreased productivity, increased development and hardware costs, and damaged customer relations” (Williams and Smith, 2002).

The deeper trouble is, even if resources are allotted towards performance operations, they rarely yield a notable difference. Once a system has been implemented atop a rigid architecture, upper-layer performance optimizations are ineffective and inconsequential. They can only offer minute upgrades within individual modules that shave mere milliseconds at a time off of the total run time. Code optimization on a modular scope offers an extremely low effort to performance improvement ratio. In fact, from personal experience optimizing scripts on a module to module basis at Veeva Systems, I judge that return from efforts diminish exponentially, as graphed below.

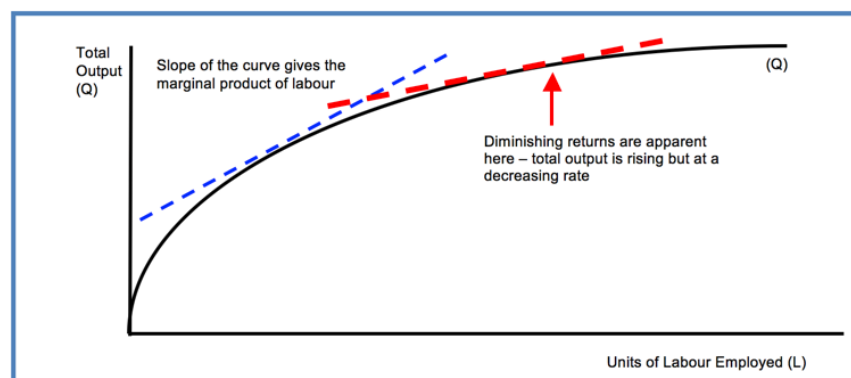


Figure 1: Diagram to show diminishing returns in performance as extra labor is employed (Riley, p1).

The difficulty in meeting software performance requirements often stems from early stage architectural design flaws. These flaws are “introduced in early development, but not discovered until late, when they are more difficult and costly to fix” (Williams and Smith, 2002). Though it may seem natural that investing more into the development phase could reduce performance issues in the long run, that is rarely the case. Due to the technical, legal and business complexity of industry-grade software, the end product is almost never an exact manifestation of the original blueprints. Thus, architectural designs seldom accurately forecast future performance issues.

Obviously, software performance issues affect software companies, be it providing software-as-a-service or software-as-a-product. These issues, if persistent in production, would impact the businesses that buy and use these software, resulting in decreased productivity, decreased revenue and increased expenses. It would also imply weaker relations between software providers and their clients.

Emphasis on Architectural Design

There exists a plethora of performance testing tools to combat the army of performance issues. As a performance engineer intern at Veeva, I learned to develop and run multiple layers of performance tests: stress tests, load tests, soak tests, component tests, end-to-end use case tests, all of which are significant in their own way. However, these tests remain unfruitful if the source architecture, especially in an embedded distributed systems application, is poorly designed. Thus, the primary focus should lie on investing extra time, resources and efforts on software architecture. Of course, no design can be flawless, and certainly not predictive. In fact, with agile movement, development should proceed adaptively as opposed to predictively. As described by the twelfth principle of the Agile Manifesto: “At regular intervals, the team reflects on how to become more

effective, then tunes and adjusts its behavior accordingly” (Beck et al., 2001). As included in the diagram below, the stage which of adapting to feedback lies in the “Start Initial Planning” phase.



Figure 2: The Agile Development Methodology Lifecycle (Softlets, p1).

I propose that at the beginning of every sprint, after release from last sprint, a formal assessment of the architecture be conducted. According to Williams and Smith, one of the best practices for performance assurance involves “performing an architecture assessment to ensure that the software architecture will support performance objectives” (Williams and Smith, 2002). Certainly, this would increase the amount of work required per sprint, which equates to less time for planning, development and testing. Nonetheless, periodical architecture performance checkpoints would greatly reduce wasted resources towards the later stages of product development, where resources can be allotted to customer satisfaction and request handling.

Emphasis on Agile Reinforcement

To further tackle performance milestones, one should look to understand the agile development methodology in depth. Agile development works in short one to four-week sprint cycles, where each sprint is structured by four ceremonies: sprint planning, daily stand-up, sprint demo and sprint retrospective (Atlassian, 2016). In respect to meeting performance objectives, there can and should be improvements to all ceremonies of the agile movement.

Performance Unit Testing

During sprint planning, product managers should consider previous sprints' progress and plan accordingly, but at least a quarter of the sprint agenda should be dedicated to performance specific efforts. Since, according to Williams and Smith, all developers and managers should be held responsible for maintaining performance standards, I propose that modular performance be a common criterion amongst unit testing (Williams and Smith, 2002). Thus, while product managers plan for performance tasks, developers implement them. Also, having unit testing cover low level performance ensures that daily scrum stand-ups bring more attention to performance. At Veeva Systems, I've spent many hours profiling features, then looking through the stack traces to determine which precise method was hogging system time and resources. Should our development team implement performance testing as a subsection of unit tests, the need for late stage profiling would diminish, thus freeing up time for more meaningful tasks.

As all performance issues arise from lower levels of source code, it is imperative that the issue is dealt with at its root. This problem demonstrates the Butterfly Effect, where a butterfly's wing flap triggers a series of events that lead to a hurricane (Ghys, 2012). An inefficient method that is called several times from multiple sessions will easily build up the server response time.

I've opened countless defect tickets at Veeva pertaining to particular methods that were performance bottlenecks. For instance, though a method itself may only take 500 milliseconds to run, if it is called 4 times when you load a page, that method alone would take up 2 seconds. Therefore, to minimize bottlenecks, I would highly propose increased resource investment into performance unit testing.

Performance Regression Testing

Moreover, I propose that at the end of every sprint, a round of regression performance tests be run in addition to the already planned general integration or regression tests. A regression test is simply the process of re-testing software that has been modified (Amman and Offut, p.215). On the Veeva Vault team, regression tests are run usually once a sprint, but regression performance tests are run at the end of the release, which is a much longer interval than it should be. Running regression performance tests allows for frequent updates on end-to-end component performance measures, which may potentially give insight on performance objectives for the following sprint. If Veeva Vault implements performance regression tests every sprint, there would be noticeably fewer performance issues when release deadlines approach.

Disadvantages to Performance Investment

On the contrary, it is imperative to consider the offset of investing resources into performance work. To reserve a significant portion of each sprint to performance efforts implies a loss of time for all other activities. Development deadlines would be tighter, which almost directly correlates to decrease in code quality and, thus, more bugs. There would be noticeably less time for writing and running tests, which means more functional bugs would eventually slip through to production.

Furthermore, from personal experience at Veeva Systems and job-hunting on Jobmine, there is currently little demand for performance engineers in the software industry. This indirectly decreases the number of qualified software performance engineers, which complicates the hiring process. More resources would be used up in finding developers that have also had sufficient training or knowledge of software performance.

Conclusion

The search for the perfect software development cycle is only a sub-task of a largely sought after answer: how does one create perfect software? There does not exist an optimal software development schedule. In fact, there may never be one. To offset the imperfect code that software developers inevitably write, significant resources and efforts are invested into quality assurance; armies of test cases, days spent re-planning sprints, and endless knowledge acquisition sessions. Amongst the vast field of quality assurance, software performance is a subject most difficult to master. To optimize corporate performance engineering, I have described in this report a useful list of considerations and practices. Specifically, during the planning and design phase, there should be a heavy focus on architecture, with thorough technical attention for scalability and responsiveness. I also stress that throughout the sprint cycle, performance-focused sprint planning be carried out, performance unit testing and performance regression testing are implemented and that sprint retrospection involve performance progress. Following the reinforced agile development cycle proposed in this report will ultimately lead to improvements in software performance. The contents of this report are effective software performance optimization tools. In fact, it is much like an innovative ski climbing utility. With time and effort, but it will get you to the flag up the hill with a promise that you won't slip. Optimizing software performance has proven itself again and again to be a worthy challenge, but that should stop no one from attempting it.

Recommendations

From a four-month co-op term at Veeva Systems as an intern performance engineer on the Vault Team, I make a few recommendations to streamline software performance efforts. Firstly, I recommend that all project owners and scrum leaders review and evaluate the present state of the system infrastructure at least once per sprint, preferably towards the beginning. I also heavily advise that Veeva Vault push for performance unit testing, so that developers can take responsibility for the performance of their code. Lastly, I urge that specialized performance engineers run performance regression tests at the end of each sprint to actively track the performance progress of the product. These recommendations apply not only to Veeva Systems, but to all software companies in general, as I firmly believe in the potential of a performance-based agile development cycle.

References

- "A Brief Introduction to Scrum." Atlassian. N.p., 2016. Web. 09 Aug. 2016. <<https://www.atlassian.com/agile/scrum>>.
- Ammann, Paul, and Jeff Offutt. Introduction to Software Testing. New York: Cambridge UP, 2008. Print.
- Beck, Kent, Mike Beedle, Arie Van Bennekum, Alistair Cockburn, Ward Cunningham, Martin Fowler, James Grenning, Jim Highsmith, Andrew Hunt, Ron Jeffries, Jon Kern, Brian Marick, Robert C. Martin, Steve Mellor, Ken Schwaber, Jeff Sutherland, and Dave Thomas. "Manifesto for Agile Software Development." Manifesto for Agile Software Development. N.p., 2001. Web. 02 Aug. 2016. <<http://www.agilemanifesto.org/>>.
- Ghys, Etienne. The Butterfly Effect (n.d.): n. pag. ENS DE LYON. 8 July 2012. Web. 3 Aug. 2016. <<http://perso.ens-lyon.fr/ghys/articles/butterflyeffect.pdf>>.
- Hein, Piet. Grooks. N.p.: n.p., 1996. Print.
- "Softlets : Embrace the Future." softlets. Softlets India, 2016. Web. 22 Aug. 2016. <<http://www.thesoftlets.com/approach.php>>.
- Williams, Lloyd G., Ph.D., and Connie U. Smith, Ph.D. "Software Performance Engineering." Best Practices in Software Measurement (2003): n. pag. Www.perfeng.com. Performance Engineering Services, 2003. Web. 2 Aug. 2016. <<http://www.perfeng.com/papers/bestprac.pdf>>.
- Williams, Lloyd G., Ph.D., and Connie U. Smith, Ph.D. "Solving Common Java EE Performance Problems." 5 Steps to Solving Software Performance Problems (2002): 351-72. Perfeng.com. L&S Computer Technology Inc., 2002. Web. 30 July 2016. <<http://www.perfeng.com/papers/step5.pdf>>.
- Woodside, Murray, Greg Franks, and Dorina C. Petriu. Sce.carleton.ca. Carleton University, 8 Feb. 2007. Web. 2 Aug. 2016. <<ftp://ftp.sce.carleton.ca/pub/cmw/07/woodside-performance.pdf>>.